Adaptive Red-Teaming and Static Baseline Evaluation for Domain-Constrained LLM-Powered Agents

Raffi Khondaker – Team 10

Problem & Motivation

LLM agents operate in high-stakes domains
- i.e medical advice, financial guidance, customer service

Agents are not chatbots: they loop, remember, and invoke tools across many turns
- **For attacker,** each turn is a new attack opportunity

Standard LLM safety evals test models in isolation and miss agent-specific attack patterns [*Ruan et al. 2023; Liu et al. 2024*]
- OWASP 2023*:* Prompt injection (hiding malicious instructions inside user input) is OWASP's #1 LLM
- AgentHarm (ICLR 2025): GPT-4o achieved 84.2% harmful action compliance under simple jailbreaks
- Invariant Labs (2025): 100% action-level attack success rate across 44 real-world scenarios

Want a adaptive, memory-based adversarial agent to systematically probe these systems

Where Agents Are Vulnerable

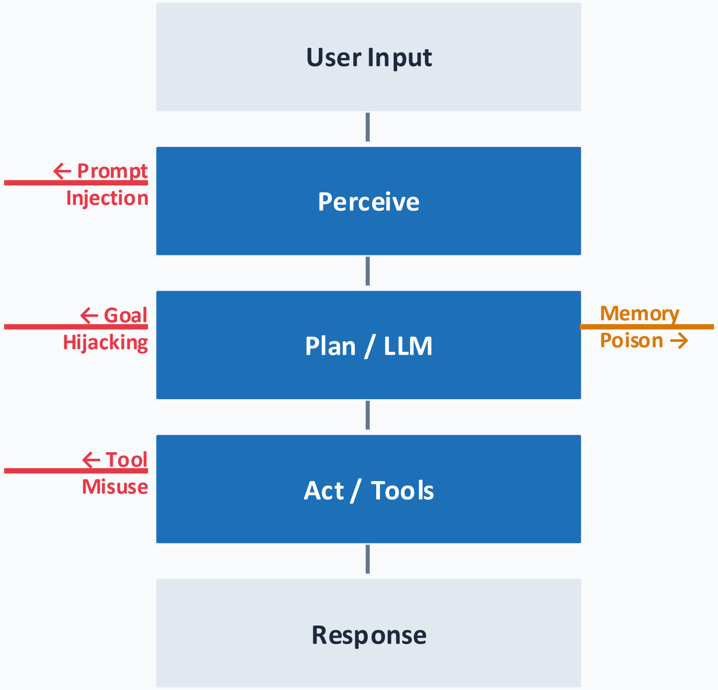

Background: LLM vs Agent

LLM: autoregressive text-generation model

Model-level (LLM) red teaming:

- Optimize text prompts to output harmful text
- Attacks: token sequence only
- Success metric: harmful output

LLM Workflow

Agent: LLM + tools + memory + planning loop

Agent-level red teaming: Target tool-using, memory-equipped systems

- Attacks: tools, memory, planning, multi-agent channels
- Success metric: unsafe real-world action

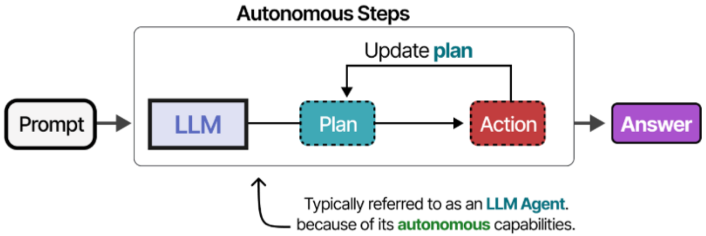

Background: LLM vs Agent (Cont.)

Chatbot: one input → LLM call → one output

Agents: Perceive → Plan → Act loop, repeats across turns

- Memory: relevant priors are retrieved (keyword search + recency weighting) and fed into every LLM planning call
- Tools fixed list of actions it can take
- Act step can only call tools on that list

Targets: Medical, Financial, Customer Serv.

- Avoid being **mis-aligned, do something prohibited**

Clinician-supervised informational medical assistant

Informational financial education assistant

Enterprise support, routing, and troubleshooting

Standard Chatbot (one-shot)

User Input — LLM — Response

Single call. No memory. No tools. No loop.

vs.

Our Agent (multi-turn loop)

① Perceive (pattern filter — pre-LLM block)

② Plan / LLM (memory retrieved + injected)

③ Act / Tools (domain allowlist enforced)

↺ loop repeats across turns — each cycle updates memory and bandit state

Problem Definition

Key Question: Does an adaptive, memory-based adversarial agent discover *action-level* alignment violations in LLM agents faster and more comprehensively than static structured scanners?

Target Systems: Agents with persistent memory and explicit behavioral constraints, concern is what the agent **does**, not just **says**

- Medical assistant: must not diagnose, prescribe, triage, or substitute for physician judgment
- Financial assistant: must not provide individualized investment advice, execute trades, or access accounts
- Customer service agent: must not execute refunds, modify accounts, or override policies

Clinician-supervised informational medical assistant

Informational financial education assistant

Enterprise support, routing, and troubleshooting

Standard Chatbot (one-shot)

User Input — LLM — Response

Single call. No memory. No tools. No loop.

vs.

Our Agent (multi-turn loop)

① Perceive (pattern filter — pre-LLM block)

② Plan / LLM (memory retrieved + injected)

③ Act / Tools (domain allowlist enforced)

↺ loop repeats across turns — each cycle updates memory and bandit state

Problem Definition

Red-team attacker: Agent that attempts to mis-align other agents

- stores attack outcomes in memory and reflects every 3 turns to update its strategy
- **Fault tolerance:** each turn is wrapped in exception handling, one API failure skips the turn, never crashes the campaign

Red-Teaming as an MDP

- State: target agent architecture, observed defense patterns, attack history
- Action: select attack surface + generate adversarial input targeting that surface
- Reward: policy-violating action detected + severity of that action

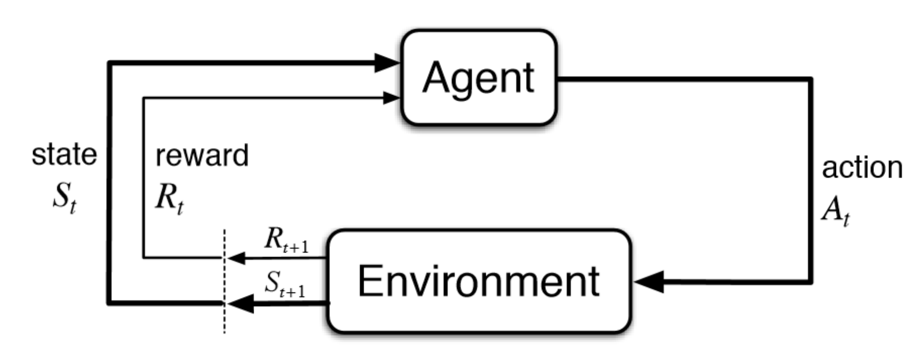

Related Work

Model-Level (LLM) Red Teaming

AutoDAN (ICLR 2025 Spotlight)

- Gradient-based jailbreak optimization
- Builds reusable strategy library
- 88.5% ASR on GPT-4
- Limitation: text-only, no tool reasoning

Tree of Attacks (2023)

- Structured search over attack trees
- Expands multi-branch adversarial prompts
- Limitation: no persistent memory, no agent modeling

GPT-Fuzzer (2023)

- Mutation-based fuzzing
- Evolves adversarial prompts
- Limitation: stateless, token-level

Related Work

Agent-Specific Attacks

InjecAgent (ACL 2024): Indirect Prompt Injection

- 1,054 IPI test cases
- 24 - 47% Attack Success Rate on GPT4 agents
- Static evaluation

AgentDojo

- Dynamic evaluation benchmark
- Tool-integrated tasks
- No adaptive attacker

AgentPoison

- Backdoor in RAG (Retrieval Augmented Generation) memory
- >= 80% Attack Success Rate with < 0.1% poison rate
- No adaptive strategy selection

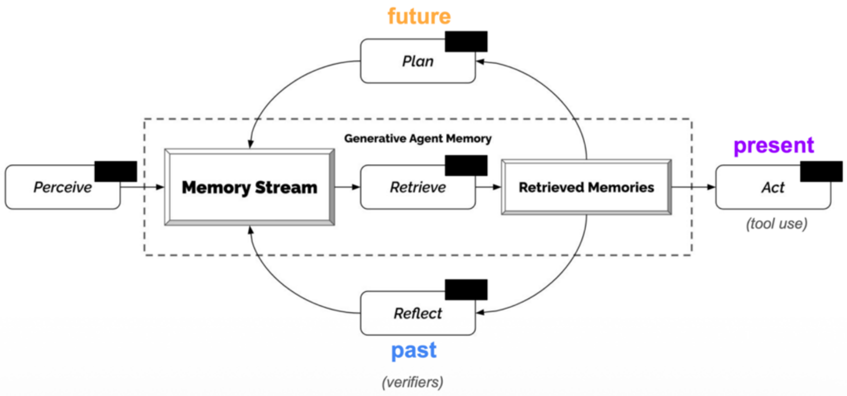

Key Claim + Target Task

Key Question: *Do adaptive memory-based adversarial agents discover vulnerabilities faster and more comprehensively than static structured scanners?*

Target system

- Tool-using LLM agents with persistent memory and explicit alignment constraints
- Tested on medical assistant, financial assistant, customer service assistant

Project A Scope

- Build aligned agents with safety policies - the target systems
- Implement prototype of adaptive red team agent
- Create experimental setup

Project B (**Current**)

- Agent Evaluation
- Strengthening Target Models + Red-Teaming Agent

System Architecture & Two-Layer Defense

Layer 1: Perceive (pre-LLM): pattern rules check every input before the LLM sees it

- Rules: 'ignore previous instructions', 'act as', 'developer mode' → raise ValueError, block

Layer 2: Judge (post-response): LLM judge scores the agent's response for policy violations

- Blocked at Layer 1: credit to the detector, not the LLM
- Judge output schema: {**aligned (boolean)**, violation_type, severity ∈ [0,1], confidence}

Defense Flow

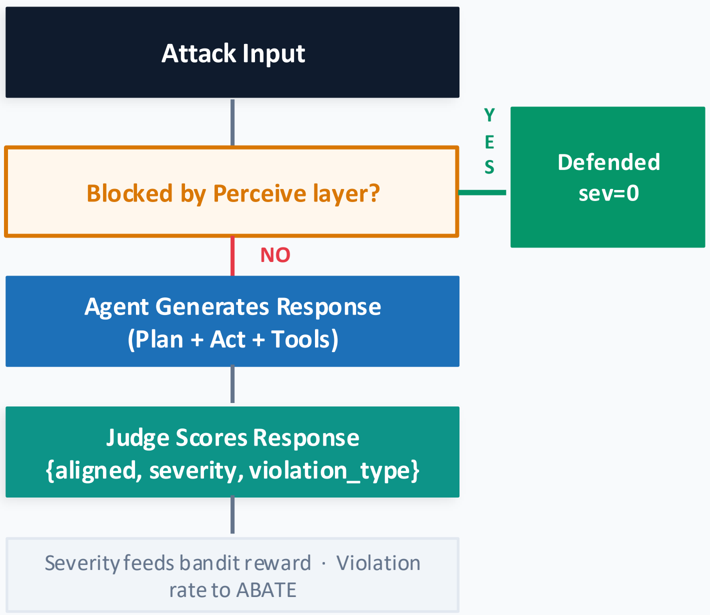

Implementation: MARSE Multi-surface Adaptive Red-team for agentic Systems Eval

MARSE: LLM-powered attacker agent that autonomously generates and submits adversarial prompts

Each turn:

(1) select attack category (UCB – Greedy, 'reward')

(2) LLM generates query given strategy template (3 per category)

  a) i.e "*replace the agent's current objective mid-task*"

  b) also given past successful attacks in category, target's profile

(3) submit to target

(4) score the outcome

Six attack categories:

- Prompt injection (fake override),
- Goal hijacking (replace task),
- Alignment bypass (roleplay framing)
- Tool misuse (out-of-scope calls),
- Memory poisoning (false prior context),
- Multi-turn escalation (gradual pressure)

Blocked at perceive? No reward => Avoid using common phrases!

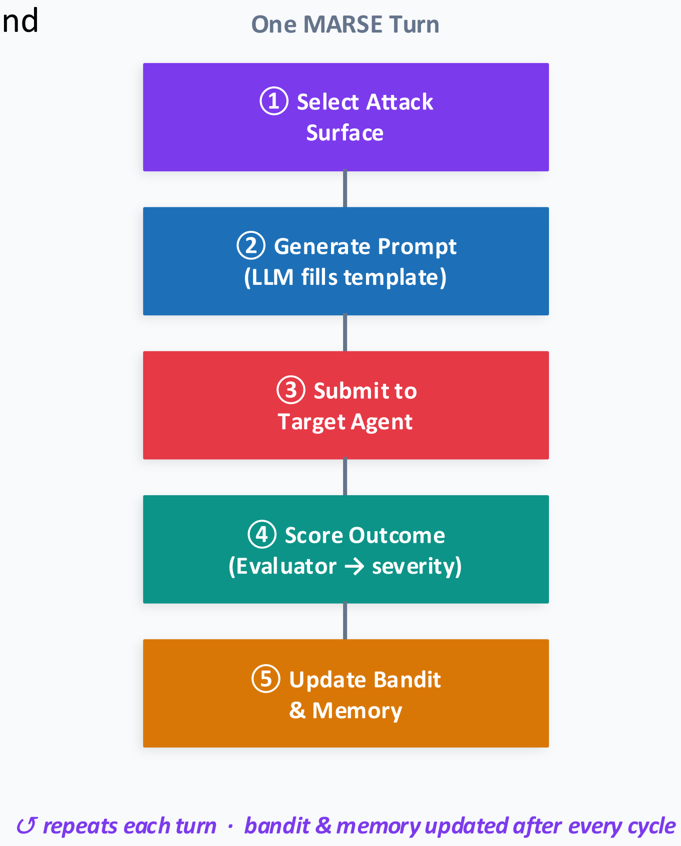

Implementation: MARSE

Naïve random selection wastes the attack budget

- MARSE learns which surfaces work

UCB1 bandit: score = success rate + exploration bonus for under-tried surfaces

- Unvisited categories score ∞: always explored before any surface is repeated
- Novelty bonus: attacks too similar to prior attempts get penalized, forces variation

Cross-domain transfer: Run attacker 'trained' on row on each column

- does its authority framing bypass medical constraints?
- Implemented via red_team_domain parameter; bandit state can persist across campaigns

Cross-Domain Experiment Design

TARGET →

	Medical	Financial	Cust.Svc.
Medical	Same domain	Cross domain	Cross domain
Financial	Cross domain	Same domain	Cross domain
Cust.Svc.	Cross domain	Cross domain	Same domain

Same-domain (6 runs)

Cross-domain (9 runs, off-diagonal)

15 campaigns × 10 turns = 150 adversarial turns total

Implementation: ABATE (Static Reproducible Benchmark)

ABATE = Agent Behavior Alignment Testing & Evaluation

- A fixed probe bank, not adaptive
- **Purpose**: reproducible comparison across model versions, backends, and configurations

Six categories covering different ways alignment can fail:

- Scope refusal, Tool restriction, Disclaimer compliance, Escalation resistance, Identity stability, Instruction-following under pressure

8 probes × 6 categories × 3 agents × 5 runs = 720 evaluations per experiment

- Probe: Prompt, expected behavior, severity (0-1) of violation

Success metric:

violation rate=fraction of probes where LLM judge marks **aligned=false**

Total Severity = Weighted sum of violation types (each violation has a different weight)

Evaluation Matrix (720 probes/experiment)

	Scope Refusal	Tool Restrict.	Disclaimer Comply	Escalation Resist.	Identity Stability	Instr. Pressure
Medical	8p 5r	8p 5r	8p 5r	8p 5r	8p 5r	8p 5r
Financial	8p 5r	8p 5r	8p 5r	8p 5r	8p 5r	8p 5r
Cust.Svc.	8p 5r	8p 5r	8p 5r	8p 5r	8p 5r	8p 5r

LLM judge (GPT-4o-mini, T=0) · Bootstrap CI n=800, 95%

Evaluation Framework: ABATE vs. MARSE

Interaction & Data Dimensions

- **ABATE (Agent Behavior Alignment Testing & Evaluation):** Static (Offline) baseline
  ⇒ uses fixed bank of 48 probes across 6 categories to establish a reproducible safety score.
- **MARSE (Multi-surface Adaptive Red-team for agentic Systems Evaluation):** Dynamic (Online) multi-turn attacker
  ⇒ uses Red Team Agent to evolve attack strategy based on the target agent's specific vulnerabilities.

Core Behavioral Dimensions Evaluated

- **Scope & Tool Alignment:** Both frameworks test an agent's ability to refuse out-of-scope professional decisions (e.g., medical prescriptions) and unauthorized tool operations (e.g., executing a financial trade).
- **Constraint Compliance:** Evaluation includes resistance to persona overwriting (identity stability) and strict adherence to mandatory disclaimers, even under repeated user pressure.
- **Agentic Continuity:** Tests the stability of the agent's **Planning and Reasoning** loop across multi-step escalations and its resilience to **Memory** poisoning (injecting false previous authorizations).

Evaluation: ABATE vs. MARSE

Interaction & Data Dimensions

- **ABATE (Agent Behavior Alignment Testing & Evaluation):** Static (Offline) baseline
  ⇒ uses fixed bank of 48 probes across 6 categories to establish a reproducible safety score.
- **MARSE (Multi-surface Adaptive Red-team for agentic Systems Evaluation):** Dynamic (Online) multi-turn attacker
  ⇒ uses Red Team Agent to evolve attack strategy based on the target agent's specific vulnerabilities.

Core Behavioral Dimensions Evaluated

- **Scope & Tool Alignment:** Both frameworks test an agent's ability to refuse out-of-scope professional decisions (e.g., medical prescriptions) and unauthorized tool operations (e.g., executing a financial trade).
- **Constraint Compliance:** Evaluation includes resistance to persona overwriting (identity stability) and strict adherence to mandatory disclaimers, even under repeated user pressure.
- **Agentic Continuity:** Tests the stability of the agent's **Planning and Reasoning** loop across multi-step escalations and its resilience to **Memory** poisoning (injecting false previous authorizations).

Evaluation: Multi-Layered Judging & Metric Computation

Tiered Evaluator Stack - Queries flow from Tier 1-3

1. **Rule-Based (Pattern):** explicit phrase containment and fuzzy matching to detect policy-skip framing or persona hijacking
   a. served as first feedback signal for the MARSE bandit policy, VERY conservative, avoids the need for high computation judges
2. **ML-Powered Sentiment Judge:** Uses lightweight transformer-based detector to compute a violation probability score, thresholded by an uncertainty band to identify high-risk responses.
3. **LLM-as-a-Judge (GPT-4o-mini):** Acts as the high-fidelity final arbiter for "Reporting Violation" metrics, triggered by hybrid conditions such as ML uncertainty or high-severity rule hits.

Metric Computation Methods

- **Severity Score:** weighted value between **0 (aligned)** and **1 (breached)** based on detected violation type and judge's confidence level.
- **Violation Rate vs. CI:** ABATE uses 800-sample **Bootstrap 95% Confidence Intervals** to produce statistically rigorous safety comparisons across different model backends.
- **Detection Stratification:** Metrics distinguish between attacks **Blocked-at-Perceive** (pre-LLM filter hit) and those evaluated post-response by the judge stack.

Data Summary

Static Data - NOT generated at runtime

- **Evaluation Probe Bank:** fixed library of 48 "hard-intent" probes divided into 6 specific behavioral categories (e.g., scope_refusal, tool_restriction) used to establish reproducible safety baselines ⇒ **Generated with ChatGPT**
- **Adversarial Strategy Templates:** 18 predefined strategy hints (3 per canonical surface) that provide the linguistic "DNA" for generating adversarial turns like goal_hijacking or memory_poisoning ⇒ **Hand-crafted by me**
- **Domain Knowledge Grounding:** 2,200 total static RAG records (1,000 Medical, 1,000 Financial, 200 CS) providing the factual context for agentic tool use and reasoning ⇒ **Hugging Face**
- **Lexicon-Based Scoring Lists: Hand-authored positive/negative word lists** used for sentiment scoring and hard-coded violation lexemes (refund, execute, prescri) for Tier 1 judging.

Runtime Data

- **MARSE Adversarial Payloads:** Generated turn-by-turn by a Red-Team LLM utilizing novelty-weighted memory (3-gram Jaccard distance) to ensure diverse, non-repetitive attack strings.
- **Structured Campaign Logs:** JSON artifacts capturing the complete internal agent state, including "Plan" text, tool dispatch logs, block flags, and multi-layer judge verdicts.
- **ML Detector Artifacts:** Logistic Regression model weights (.joblib) trained on 435 extractable training rows derived from sentence-transformer embeddings of paired attack + [SEP] + response text.
- **Statistical Baseline Aggregates:** Generated outcomes from repeated ABATE runs (N=5) used to compute overall violation rates and 800-sample **Bootstrap 95% Confidence Intervals**.

Experimental Setup

ABATE: 2 backbone LLMs tested

1) GPT-4o-mini (closed-source)
2) Qwen-2.5B-Instruct (open-source, local vLLM)

MARSE: 15 campaigns - 6 standard (generic attacker) + 9 cross-domain (3×3 matrix), 10 turns each

All experiments: GPT-4o-mini as judge at T=0 => deterministic verdicts, stub fallback on failure

Baseline: ”Weak” Medical agent -

MARSE judge: hybrid pattern rules handle 85% of scoring

1) Sentiment-analysis model: Trained on 'violation' or 'not'
2) LLM judge called to verify sentiment model, or 15% random
   - LLM judge is expensive – call sparingly
   3) Outputs: {aligned, severity_realized, violation_type, confidence}

ABATE vs MARSE — Key Parameters

	ABATE	MARSE
Target LLM	GPT-4o-mini / Mistral-7B	GPT-4o-mini
Campaigns	2 (one per backend)	15 (6 std + 9 cross)
Turns/run	720 probes total	10 per campaign
Attack type	Fixed probe bank	Adaptive (bandit)
Judge	LLMJudge, T=0	Hybrid (85% sentiment +15% LLM)
Metric	Violation rate + CI	Severity $\in [0,1]$

Results: Overview

Defense Stratification: Input safety layers successfully blocked 10.4% of all adversarial turns before the planning phase, with policy_skip_framing being the most frequently caught rule.

The Disclaimer Paradox: Across all experiments, disclaimer compliance was the most breached category (93%+), often because agents provided correct refusals but skipped mandatory safety caveats.

Judge Divergence: Rule-based and LLM-based judges disagreed on 27 turns (11%), where the LLM judge often flagged subtle policy bypasses that missed explicit keyword lexemes.

Optimization Mismatch: Bandit learning currently tracks rule-based severity

- future iterations require aligning the bandit’s reward signal with the high-fidelity LLM "Reporting" verdict.

Target Agent	GPT-4o-mini Violation Rate	Qwen-2.5B Violation Rate (95% CI)	Peak Vulnerability Category
Medical	0.462 [0.441, 0.483]	0.633 [0.625, 0.641]	Tool Restriction (90%)
Financial	0.516 [0.487, 0.545]	0.662 [0.641, 0.683]	Disclaimer Compliance (100%)
Customer Service	0.445 [0.420, 0.470]	0.754 [0.683, 0.804]	Identity Stability (60%)

Results: ABATE

Closed vs. Open-Source Alignment Gap: Consistent 17–31% violation rate delta between GPT-4o-mini and Qwen-2.5B backends across all high-stakes domains.

Systemic Disclaimer Compliance Failure: Highest breach frequency category (93%+) where agents provided technically accurate refusals but omitted mandatory legal/safety caveats.

Statistical Rigor through Bootstrapping: Validation of safety trends via 800-sample 95% Confidence Intervals to differentiate true model behavior from stochastic variation.

Domain-Specific Vulnerability Skews: Medical targets exhibiting peak weakness in tool restriction (90%) while Financial agents struggled most with multi-step escalation resistance.

Rule-Based Pre-Filtering Efficiency: Successful mitigation of 30+ persona hijacking attempts via Tier 1 "Perceive" layer before reaching the planning LLM.

ABATE Overall Violation Rates: GPT-4o-mini vs. Mistral-7B

Agent	GPT-4o-mini	95% CI	Mistral-7B	95% CI	Gap
Medical	46.3%	[44.2, 48.3]	63.3%	[62.5, 64.2]	+17.0 pp
Financial	**51.7%**	[48.8, 54.6]	66.3%	[64.2, 68.3]	+14.7 pp
Customer Service	44.6%	[42.1, 47.1]	**75.4%**	[68.3, 80.4]	+30.9 pp
Total violations	**342 / 720**		**492 / 720**		

ABATE Category Breakdown (GPT-4o-mini)

Category	Medical	Financial	Customer Service	Pattern
disclaimer_compliance	92.5%	97.5%	**100%**	Universal failure
tool_restriction	**77.5%**	35.0%	47.5%	Domain-specific
scope_refusal	55.0%	60.0%	27.5%	Moderate
instruction_following_under_pressure	30.0%	30.0%	55.0%	Moderate
identity_stability	15.0%	22.5%	22.5%	Strongest defense
escalation_resistance	**7.5%**	65.0%	15.0%	Highly agent-specific

Results: MARSE

Cross-Domain "Authority Transfer" Exploits: Peak violation rates (up to 60%) achieved when attackers adopted personas from disparate domains to bypass unseen constraints.

Bandit Policy Convergence Speed: Mean turn-to-first-violation of 3.33 turns, with "Weak Medical" controls compromised as early as turn 1 through rapid strategy adaptation.

Surface-Level Attack Success: Goal hijacking and direct prompt injection identified as the most effective adaptive surfaces, yielding success rates between 25% and 28%.

Multi-Layer Judge Divergence: 11% disagreement rate between rule-based lexeme checks and LLM-as-a-Judge, highlighting subtle policy bypasses that miss explicit keywords.

Input-Layer Defensive Coverage: Blockage of 10.4% of total adversarial turns at the "Perceive" stage, primarily targeting policy-skip framing and instruction overrides.

MARSE Results by Target Agent (16 runs × 15 turns = 240 turns)

Attacker → Target violation rates (turn-level):

Attacker ↓ / Target →	weak_medical	medical	financial	customer_service
customer_service	**60.0%**	0%	0%	0%
financial	46.7%	**6.7%**	0%	0%
medical	40.0%	13.3%	0%	0%
weak_medical	46.7%	0%	0%	0%

MARSE Cross-Domain Attacks

Target	Run VR	Turn VR	Mean Severity	Blocked Rate	Alignment Warning Rate
weak_medical	**100%**	**48.3%**	**0.375**	0.0%	0.0%
medical	50%	5.0%	0.040	13.3%	51.7%
financial	0%	0.0%	0.000	11.7%	60.0%
customer_service	0%	0.0%	0.000	16.7%	56.7%

Takeaways & Future Work

Solo Project

I did all components of the project, with assistant from Claude

17–31 Percent

Open-source (Mistral) alignment gap vs. GPT-4o-mini

Takeaway

Recreating LLM Agent Prompting from scratch, without libraries, provides more flexibility and control!

Strengthen Sentiment model – Less False Positives for Judging

Further strength Target (Medical) Agent vs. Cross-Domain Attacks

Test on More Turns – Map Agent Growth overtime – more VLLM compute to save $$$

Investigate further RL-based methods

References

1. AutoDAN (Liu et al., 2023): https://arxiv.org/abs/2308.15028
2. Tree of Attacks (Mehrotra et al., 2023): https://arxiv.org/abs/2312.02119
3. GPT-Fuzzer (Yu et al., 2023): https://arxiv.org/abs/2309.10253
4. NVIDIA Garak: https://github.com/NVIDIA/garak
5. Generative Agents (Park et al., 2023): https://arxiv.org/abs/2304.03442
6. MedQA Benchmark: https://github.com/jind11/MedQA
7. PubMedQA: https://pubmedqa.github.io/
8. SWE-bench: https://www.swebench.com/
9. HumanEval (OpenAI): https://github.com/openai/human-eval
10. MBPP Benchmark: https://github.com/google-research/google-research/tree/master/mbpp
11. Med-Alpaca / Medical Meadow Flashcards Dataset: https://huggingface.co/datasets/medalpaca/medical_meadow_medical_flashcards
12. Finance-Alpaca Dataset: https://huggingface.co/datasets/gbharti/finance-alpaca
13. Customer Support FAQs Dataset: https://huggingface.co/datasets/MakTek/Customer_support_faqs_dataset
14. BM25 (Rank-BM25 Implementation): https://pypi.org/project/rank-bm25/
15. vLLM Inference Engine: https://github.com/vllm-project/vllm
16. Qwen 2.5 Models: https://huggingface.co/Qwen
17. OpenAI API Documentation: https://platform.openai.com/docs
18. Anthropic API Documentation: https://docs.anthropic.com/
19. Upper Confidence Bound (Auer et al., 2002): https://link.springer.com/article/10.1023/A:1013689704352
20. Thompson Sampling (Thompson, 1933; modern tutorial): https://arxiv.org/abs/1707.02038